CANINE vs BERT on NLP Tasks

Dina EL ZEIN dina.el-zein@ens-lyon.fr
Alexandre MARTHE alexandre.marthe@ens-lyon.fr

April 5, 2022

1 Introduction

This report is based on the work of J. Clark et al [1]. Pre-trained text encoders have superseded the state of the art on many natural language processing (henceforth NLP) tasks, yet they still require a tokenization step. In this paper, the idea is to build a free tokenizer encoder and vocabulary that works directly on character sequences, thus CANINE[1]. The code and checkpoints are available on GitHub at https://github.com/google-research/language/tree/master/language/canine.

Through this report, an overview of the model introduced in this article by J. Clark et al [1] is presented. Experiments on new datasets using CANINE are proposed with a direct comparison with mBERT [2]. The code to replicate the experiments is released at https://github.com/dinalzein/CANINE-VS-BERT-NLP-Tasks

2 Method

CANINE, the character level model is simply implemented by taking as input a sequence of characters instead of subwords. The implementation consists of there main parts: (1) a downsampling function DOWN, (2) a primary encoder ENCODE, (3) an upsampling function UP. Given an input sequence of character embeddings $e \in \mathbb{R}^{n \times d}$ with length n and dimensionality d:

$$Y_{seq} = UP(ENCODE(DOWN(e))) \tag{1}$$

such that $Y_{seq} \in \mathbb{R}^{n \times d}$ is the final representation for sequence prediction tasks while for classification tasks, the model simply takes the zeroth element of the primary encoder:

$$y_{cls} = [ENCODE(DOWN(e))]_0 \tag{2}$$

The input representation to CANINE is a sequence of integers consisting of the codepoint integer of the characters. To represent all the unicode [2] characters using only a small number of parameters, they use different hash functions to embed these codepoints integers. This way, the model can learn representation for characters during fine-tuning never seen during pre-training.

Going through the implementation parts, firstly, for the downsampling function, an encoding of the characters using a single-layer block-wise local attention tranformer has been done following by strided convolution to reduce the number of sequence positions to be similar to that of a word

---

[1]CANINE: Character Architecture with No tokenization In Neural Encoders.
[2]codepoint values are part of the Unicode Standard

piece model. Secondly, for the primary encoder, a deep transformer stack with different layers is applied to the downsampled output. Lastly, for the upsampling function, a reconstruction for the character-level output is done by concatenating the output of the original character transformer with the downsampled representation output by the deep transformer stack.

3   Data

General Language Understanding Evaluation (GLUE)[3] benchmark is a collection of natural language understanding tasks on different datasets [4]. We use the following datasets from this benchmark:

3.1   Sentiment Analysis (SST-2)

SST-2 [4] dataset [3] consists of 70042 sentences taken from movie reviews with human annotations of their sentiment. Each sample is linked with a label to tell whether or not the sentence is semantically correct. The dataset is split with 67349 sentences for training, 872 for validating, and 1821 for testing.

Metric: For evaluation, we use the accuracy metric which measures how many samples are correctly classified among all samples. The higher the value is, the better the result is.

Task: A binary classification task to predict whether the movie review is semantically correct or not.

Setup: To fine-tune on SST-2, we use a batch size of 16 for 5 epochs with learning rate $= 2e - 5$ and weight decay $= 0.01$. Then we select the best fine-tuning model on the training set and use it for evaluation.

3.2   The Stanford Question Answering Dataset (SQuAD)

The Squad[5] dataset is an English-language Question Answering dataset of 98169 questions generated from 536 Wikipédia articles. Each sample consists in a question, a context, and an answer. There are 87599 training samples and 10570 validation ones.

Metric: For evaluation, we use the F1-score, which is the measure of overlap between the predicted answer and the correct answer, and is defined as the harmonic mean of the precision and recall.

Task: A question answering task where the answer is to be found among a specific context.

Setup: We pre-process the data to add the labels for the start and end position of the answer in the context. Afterwards, we fine-tune the model on the training dataset with the following parameters: a batch size of 4, 3 epochs, a leaning rate of $2e - 3e$, and weight decay of 0.01. For the evaluation, we pre-process the data again to match logits and the characters in the context.

---

[3]https://gluebenchmark.com/leaderboard
[4]https://nlp.stanford.edu/sentiment/index.html
[5]https://rajpurkar.github.io/SQuAD-explorer/

Model	Input	Training Loss	Evaluation Loss	Accuracy
BERT	Subwords	0.060	0.334	0.927
CANINE-S	Subwords	0.157	0.625	0.851
CANINE-C	Chars	0.169	0.572	0.856

Table 1: Comparison between BERT and CANINE on SST-2

Model	Input	Training Loss	Evaluation Loss	Exact Match	F1-score
BERT	Subwords	0.773	1.157	76.746	85.134
CANINE-S	Subwords	0.629	1.461	72.526	82.183
CANINE-C	Chars	0.664	1.355	72.375	82.300

Table 2: Comparison between BERT and CANINE on SQuAD

3.3 The Corpus of Linguistic Acceptability (CoLA)

CoLA[6] dataset [5] consists of 10657 English sentences taken from 23 different linguistics publications. Each sample in this dataset, which is a sequence of words, is associated with a label to tell whether this sentence is grammatically acceptable or not. From this dataset, 8551 sentences are used for training, 1043 sentences for validating, and 1063 for testing.

Metric: For evaluation, we use Matthews correlation coefficient metric which measures the association for two binary values and ranges from -1 to 1: 1 indicating a perfect match, -1 indicating a perfect disagreement, and 0 indicates an uninformed guessing (simailr to random) between the predicted label and the observed one.

Task: The task on this dataset is a binary classification task to predict whether an English sentence is grammatically correct or not.

Setup: To fine-tune on CoLA, we use the same setup as for SST-2.

4 Results

In this section we present CANINE and BERT fine-tuning results on the NLP tasks presented in Section 3.

To fine-tune using CANINE, we use checkpoints from the two models: **CANINE-C** [7] pre-trained with autoregressive character loss and **CANINE-S** [8] pre-trained with subword loss. Both models were pre-trained with 12-layers, 768-hidden, 12-heads, and $121M$ parameter. For BERT, we use the one trained on lower-cased English text that consists of 12-layers, 768-hidden, 12-heads, and $110M$ parameters.

Table 1 represents the results for a comparison between BERT and CANINE for SST-2 classification task. CANINE-C improves over CANINE-S by 0.005 negligible accuracy. BERT outperforms both CANINE-S and CANINE-C by $+0.076$ and $+0.071$ accuracy respectively.

---

[6] https://nyu-mll.github.io/CoLA/
[7] https://storage.googleapis.com/caninemodels/canine-c.zip
[8] https://storage.googleapis.com/caninemodels/canine-s.zip

Model	Input	Training Loss	Evaluation Loss	Matthews Correlation
BERT	Subwords	0.149	0.766	0.565
CANINE-S	Subwords	0.615	0.638	0
CANINE-C	Chars	0.567	0.668	0.064

Table 3: Comparison between BERT and CANINE on CoLA

Results presented in Table 2 represents the performance of fine-tuning CANINE and BERT on question answering task, SQuAD. BERT outperforms both CANINE-C and CANINE-S by +2.834 F1 and +2.951F1 respectively.

Table 3 represents the results for a comparison between BERT and CANINE for COLA classification task. Fine-tuning using CANINE-S (CANINE with the subword loss) shows an inefficient performance as Matthews Correlation = 0 means that prediction is randomized. Using CANINE-C (CANINE with the character loss), the Matthes Correlation = 0.064, slightly better than CANINE-S but still not comparable to BERT. CANINE is not good for this task at all.

5 Conclusion

We presented a comparison between the state-of-the-art NLP models CANINE and BERT on 3 different NLP tasks: SQuAD, SST-2, and CoLA. In fact, BERT outperforms CANINE for all the NLP tasks presented in Section 3 interms of F1 score, accuracy, and Matthews Correlation.

References

[1] Jonathan H Clark et al. "Canine: Pre-training an efficient tokenization-free encoder for language representation". In: *arXiv preprint arXiv:2103.06874* (2021).

[2] Jacob Devlin et al. "Bert: Pre-training of deep bidirectional transformers for language understanding". In: *arXiv preprint arXiv:1810.04805* (2018).

[3] Richard Socher et al. "Recursive deep models for semantic compositionality over a sentiment treebank". In: *Proceedings of the 2013 conference on empirical methods in natural language processing*. 2013, pp. 1631–1642.

[4] Alex Wang et al. "GLUE: A multi-task benchmark and analysis platform for natural language understanding". In: *arXiv preprint arXiv:1804.07461* (2018).

[5] Alex Warstadt, Amanpreet Singh, and Samuel R Bowman. "Neural Network Acceptability Judgments". In: *arXiv preprint arXiv:1805.12471* (2018).